



NATURAL LANGUAGE PROCESSING PROJECT

Group 4

Stefani Guevara

Mariko McDougall

Arathi Nair

CONTENTS

- Glue Benchmark Tasks
- ELECTRA
- XLNet
- DeBERTa
- Ensemble Model
- Experiment Results
- Conclusion

GLUE BENCHMARK

Single Sentence Tasks

Cola - Grammatical Correctness

SST-2 - Sentiment Analysis

Similarity & Paraphrase Tasks

MRPC - Paraphrase detection of two sentences

QQP - Paraphrase detection of two questions

STS-B - Sentence similarity

Inference Tasks

MNLI - Sentences match/mismatch

QNLI - Question & Answer pairing

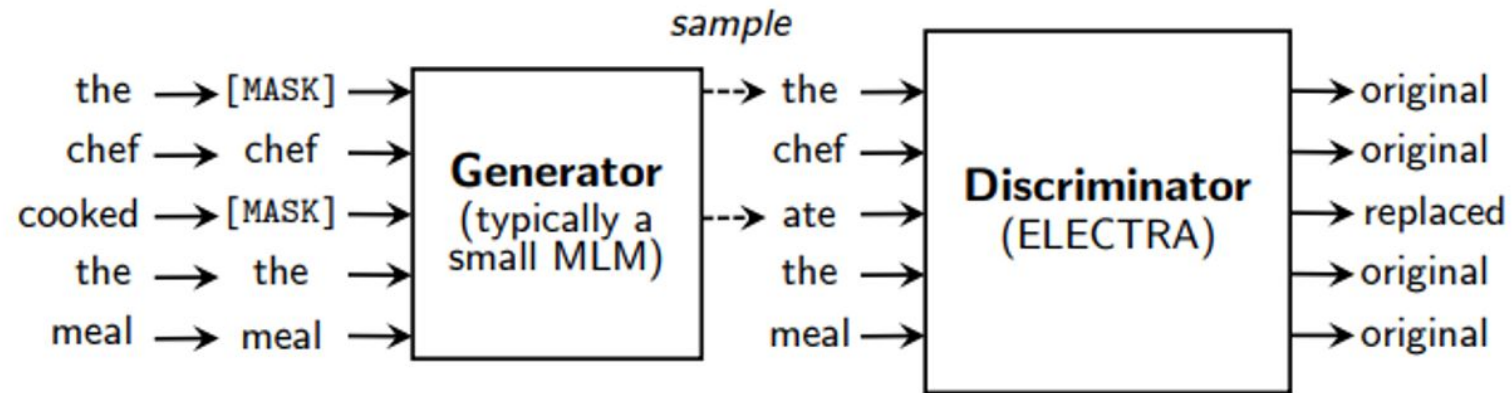
RTE - Sentences match/mismatch

WNLI - Sentences match/mismatch with pronoun substitution

ELECTRA

- Introduced a unique pre-training approach called “*replaced token detection*”
- Predicts all the input token unlike its predecessors that relied on MLM pre-training and predicts only 15% of the tokens
- Uses significantly less compute resources
- Results match or exceed downstream performance of a pre trained MLM
- More efficient pre-training compared to MLM
- Produces an improved comprehension of context

Model Architecture



1. MLM selects a random set of positions to mask out $m = [m_1, \dots, m_k]$
2. Generator predicts original words of the [MASK] tokens
3. Discriminator distinguishes tokens replaced by the generator
4. Model is trained to distinguish "real" input tokens vs "fake" input tokens

GENERATOR

Output probability for a token x_t with softmax layer

$$p_G(x_t|\mathbf{x}) = \exp(e(x_t)^T h_G(\mathbf{x})_t) / \sum_{x'} \exp(e(x')^T h_G(\mathbf{x})_t)$$

\mathbf{x} = sequence on input tokens

$h(\mathbf{x})$ = contextualized vector representations

e = token embeddings

Loss function:

$$\mathcal{L}_{\text{MLM}}(\mathbf{x}, \theta_G) = \mathbb{E} \left(\sum_{i \in \mathbf{m}} -\log p_G(x_i | \mathbf{x}^{\text{masked}}) \right)$$

DISCRIMINATOR

Predicts if token x_t is “real”, with a sigmoid output layer

$$D(\mathbf{x}, t) = \text{sigmoid}(w^T h_D(\mathbf{x})_t)$$

t = position of the token

h_D = hidden layers, embedding layers, attention heads

w^T = predicted output

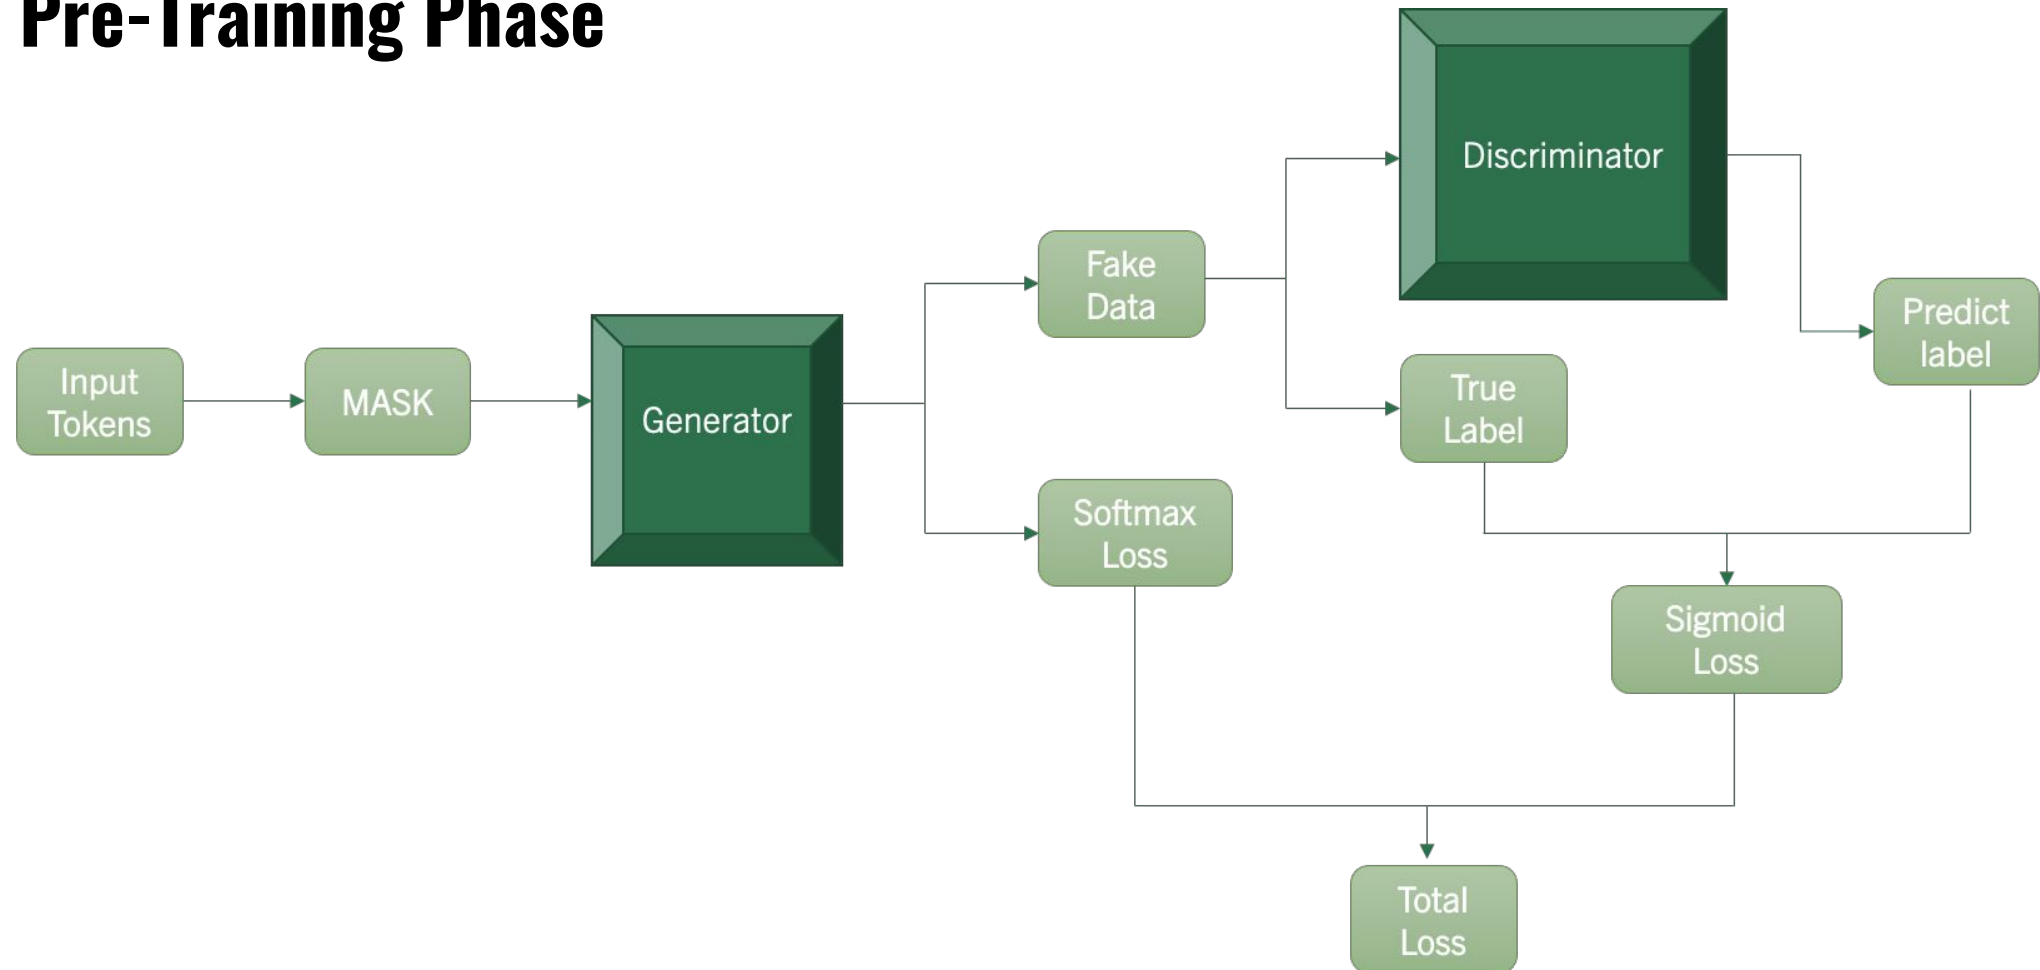
Loss function:

$$\mathcal{L}_{\text{Disc}}(\mathbf{x}, \theta_D) = \mathbb{E} \left(\sum_{t=1}^n -\mathbb{1}(x_t^{\text{corrupt}} = x_t) \log D(\mathbf{x}^{\text{corrupt}}, t) - \mathbb{1}(x_t^{\text{corrupt}} \neq x_t) \log(1 - D(\mathbf{x}^{\text{corrupt}}, t)) \right)$$

Combined Loss:

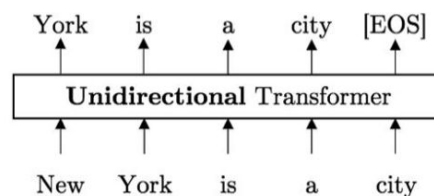
$$\min_{\theta_G, \theta_D} \sum_{\mathbf{x} \in \mathcal{X}} \mathcal{L}_{\text{MLM}}(\mathbf{x}, \theta_G) + \lambda \mathcal{L}_{\text{Disc}}(\mathbf{x}, \theta_D)$$

Pre-Training Phase



XLNET

Auto-regressive Language Modeling

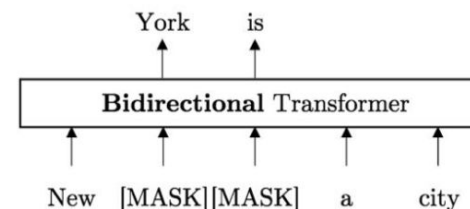


😊 Full Auto-regressive **Dependence**

😊 Free from artificial **Noise**

😞 No **Bidirectional Context**

Denoising Auto-encoding (BERT)



😞 **Independent** Predictions

😞 Artificial **Noise**: [MASK]

😊 Natural **Bidirectional Context**

XLNET

General Autoregressive Model

- Bidirectional with permutation operation by maxing joint probability

*Peter's **cat** likes yarn*

<i>Peter's cat likes yarn</i>	<i>yarn Peter's cat likes</i>
<i>Peter's cat yarn likes</i>	<i>yarn Peter's likes cat</i>
<i>Peter's likes cat yarn</i>	<i>yarn cat Peter's likes</i>

$$\mathcal{J}_{\text{BERT}} = \log p(\text{New} \mid \text{is a city}) + \log p(\text{York} \mid \text{is a city}),$$

$$\mathcal{J}_{\text{XLNet}} = \log p(\text{New} \mid \text{is a city}) + \log p(\text{York} \mid \text{New, is a city}).$$

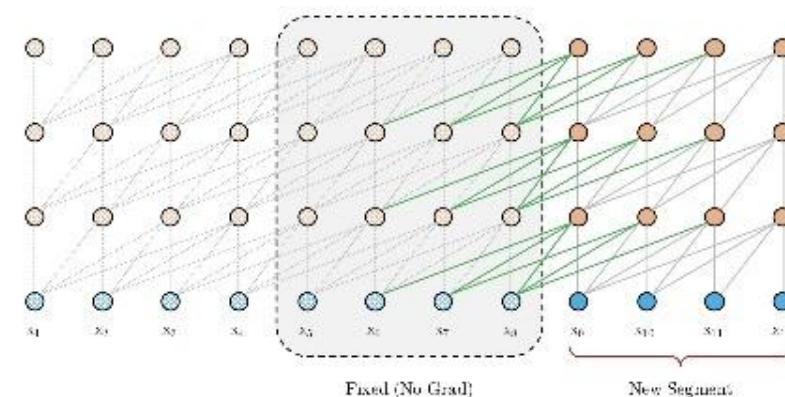
XLNET

General Autoregressive Model

- Bidirectional with permutation operation by maxing joint probability

Transformer-XL (SOTA LM) Integration

- Long-term dependencies via cache and reuse of previous hidden states



Google AI blog, [Transformer-XL: Unleashing the Potential of Attention Models](#)

XLNET

General Autoregressive Model

- Bidirectional with permutation operation by maxing joint probability

Transformer-XL (SOTA LM) Integration

- Long-term dependencies via cache and reuse of previous segment hidden states

Two-Stream Attention Mechanism

- Query: Keep positional encoding, blind to target
- Content: Gather context information (permutation)

$$g(z_t, \mathbf{x}_{\mathbf{z}_{<t}}) = \text{Attn}_\theta \left(\underbrace{\mathbf{Q} = \text{Enc}(z_t)}_{\text{Stand at } z_t}, \underbrace{\mathbf{KV} = \mathbf{h}(\mathbf{x}_{\mathbf{z}_{<t}})}_{\text{Gather info. from } \mathbf{x}_{\mathbf{z}_{<t}}} \right)$$

DeBERTa

Decoding-Enhanced **BERT** with disentangled **A**ttention

Google
BERT



DeBERTa

Decoding-Enhanced **BERT** with disentangled **A**ttention

- Disentangled Attention
 - “I love **deep learning**.” vs “We are **learning** about modeling techniques and **deep** neural nets.”
 - BERT**: vector(position embedding + word embedding) -> Attention Mask
 - DeBERTa**: position vector, word vector-> Attention Mask
 - Effect**: Attention weights of words modulated by their relative positions

DeBERTa

Decoding-Enhanced **BERT** with disentangled **A**ttention

- Disentangled Attention
 - “I love **deep learning**.” vs “We are **learning** about modeling techniques and **deep** neural nets.”
 - BERT: vector(position embedding + word embedding) -> Attention Mask
 - DeBERTa: position vector, word vector-> Attention Mask
 - Effect: Attention weights of words modulated by their relative positions
- Enhanced Decoding Mask
 - “A new **store** opened beside the new **mall**.”
 - BERT: -> Softmax layer to decode masked words
 - DeBERTa: Absolute position embeddings -> Softmax layer
 - Effect: Absolute position of words are taken into account

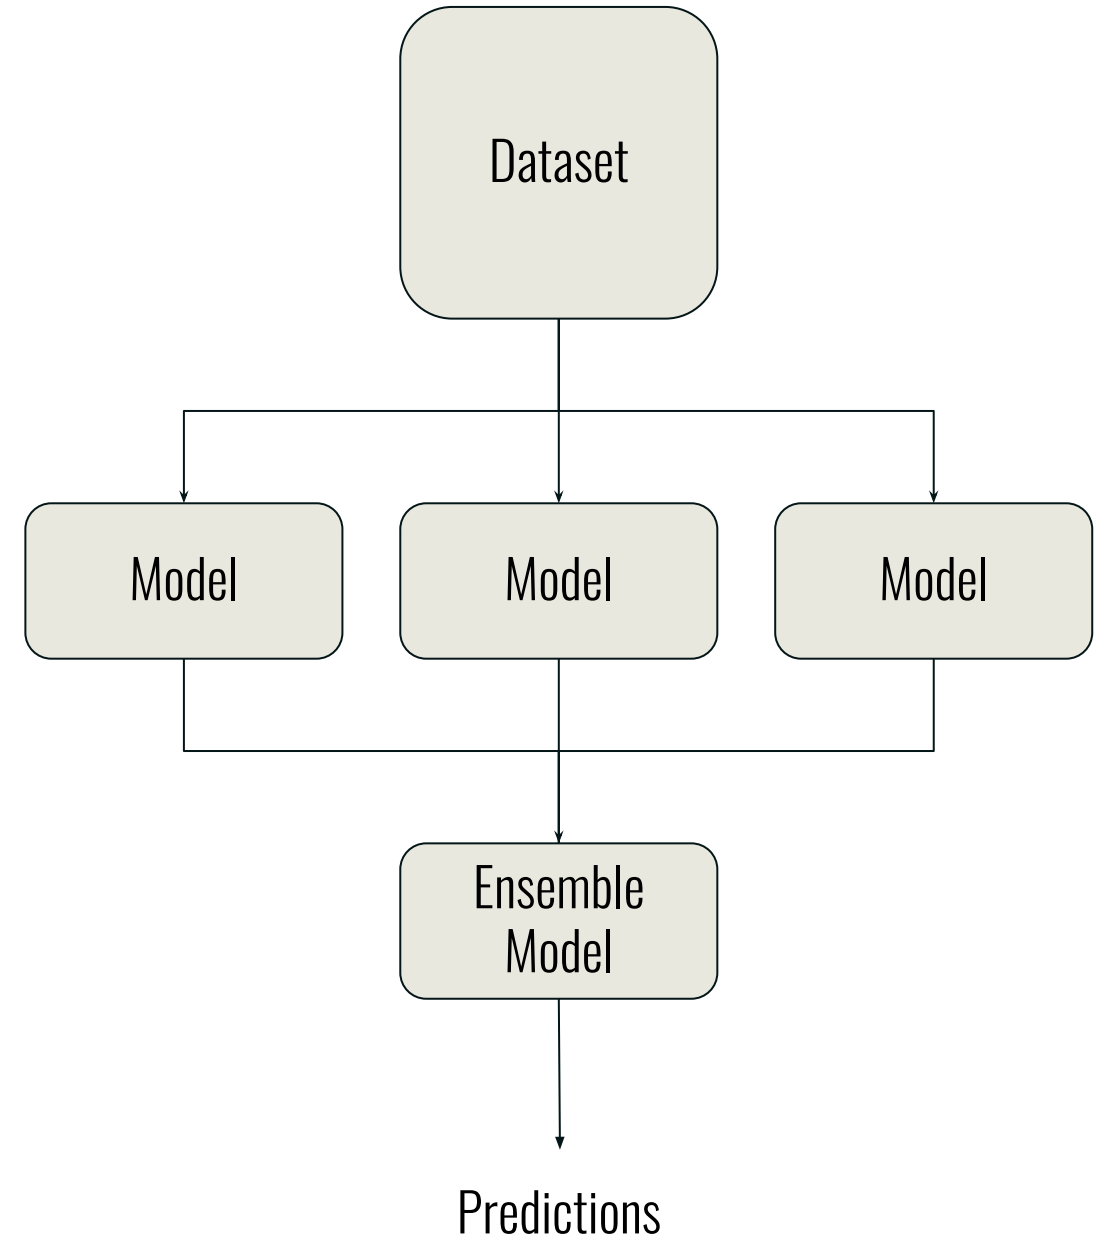
DeBERTa

Decoding-Enhanced **BERT** with disentangled **A**ttention

- Full Size Deberta V3
 - 24 layers, hidden size of 1024
 - 304M backbone parameters
 - Vocabulary of 128K tokens = 131M parameters
- Deberta V3-Small
 - 6 layers, hidden size of 768
 - 44M backbone parameters
 - Vocabulary of 128K tokens = 98M parameters

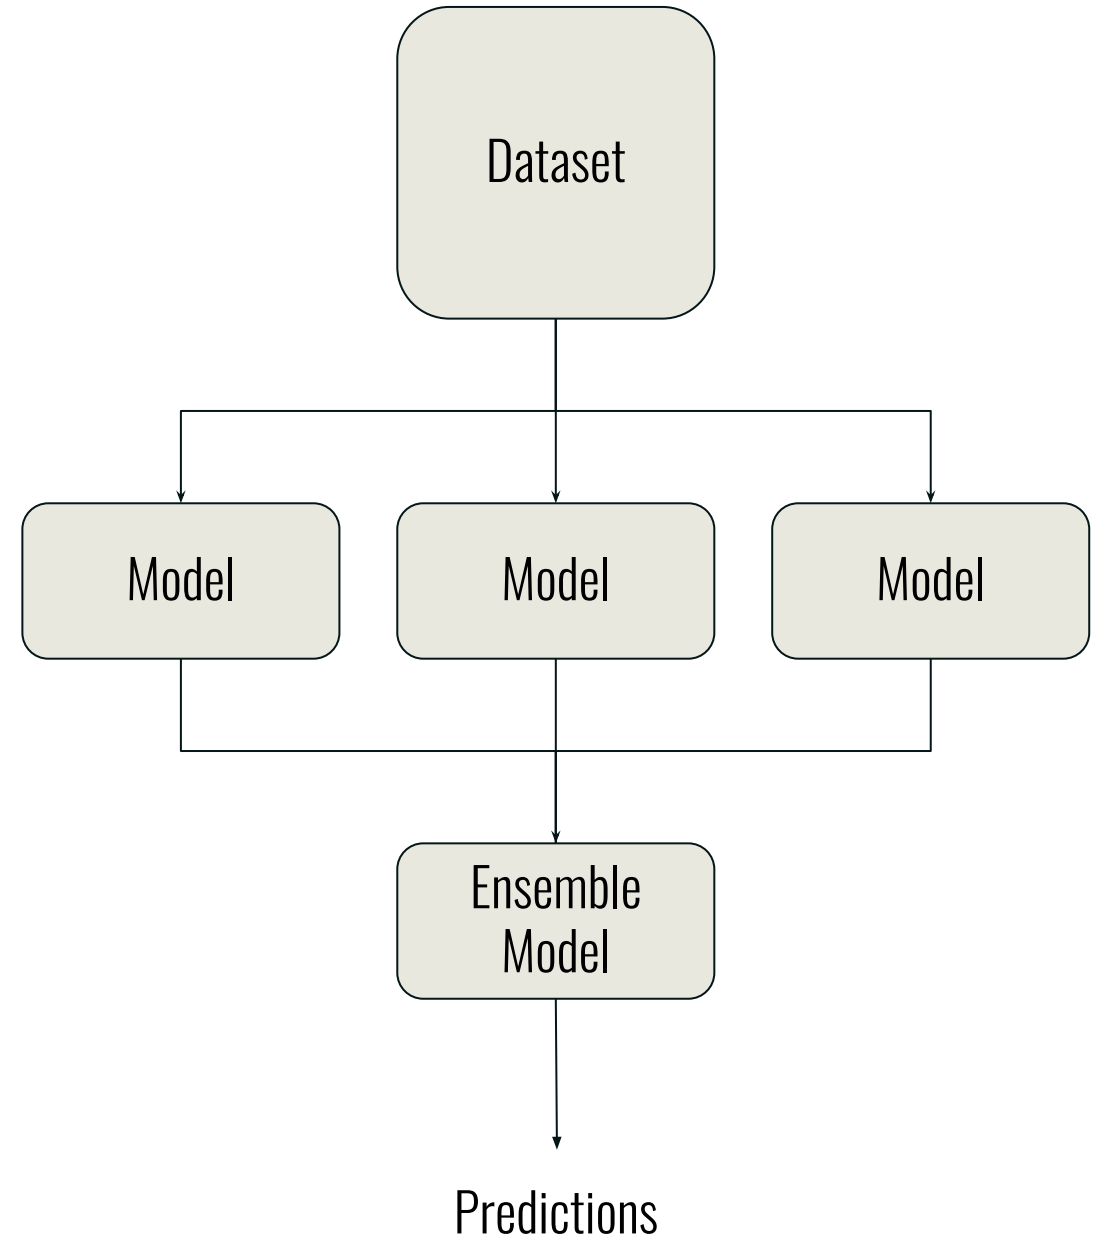
Ensemble Model

- **Ensemble learning**
Meta approach to modeling
Leverages the power of multiple predictive models

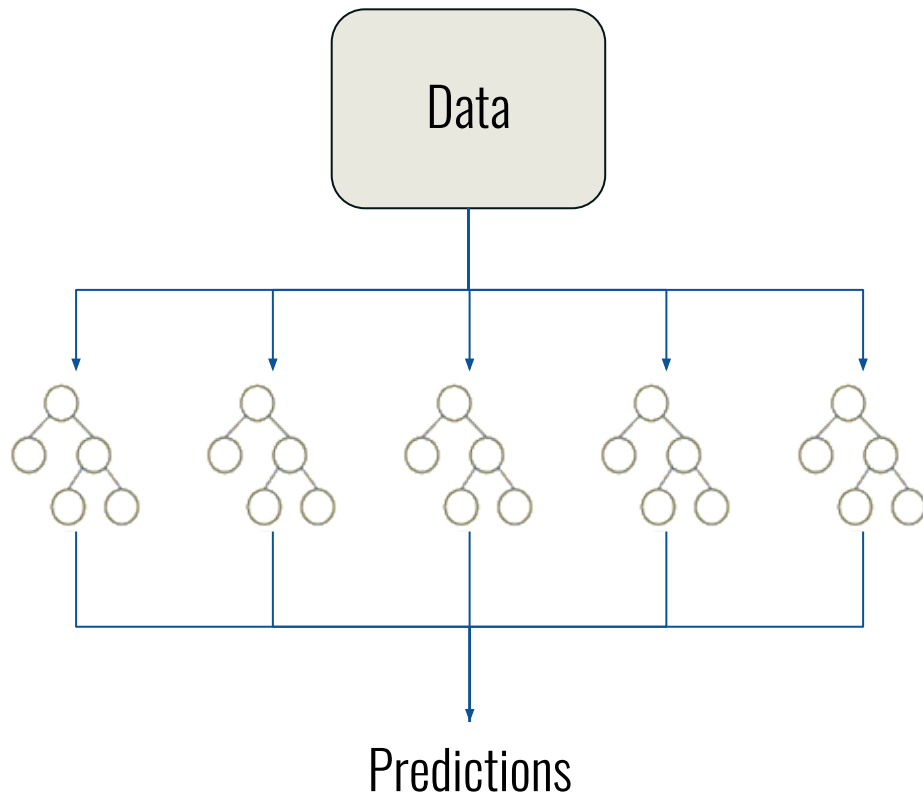


Ensemble Model

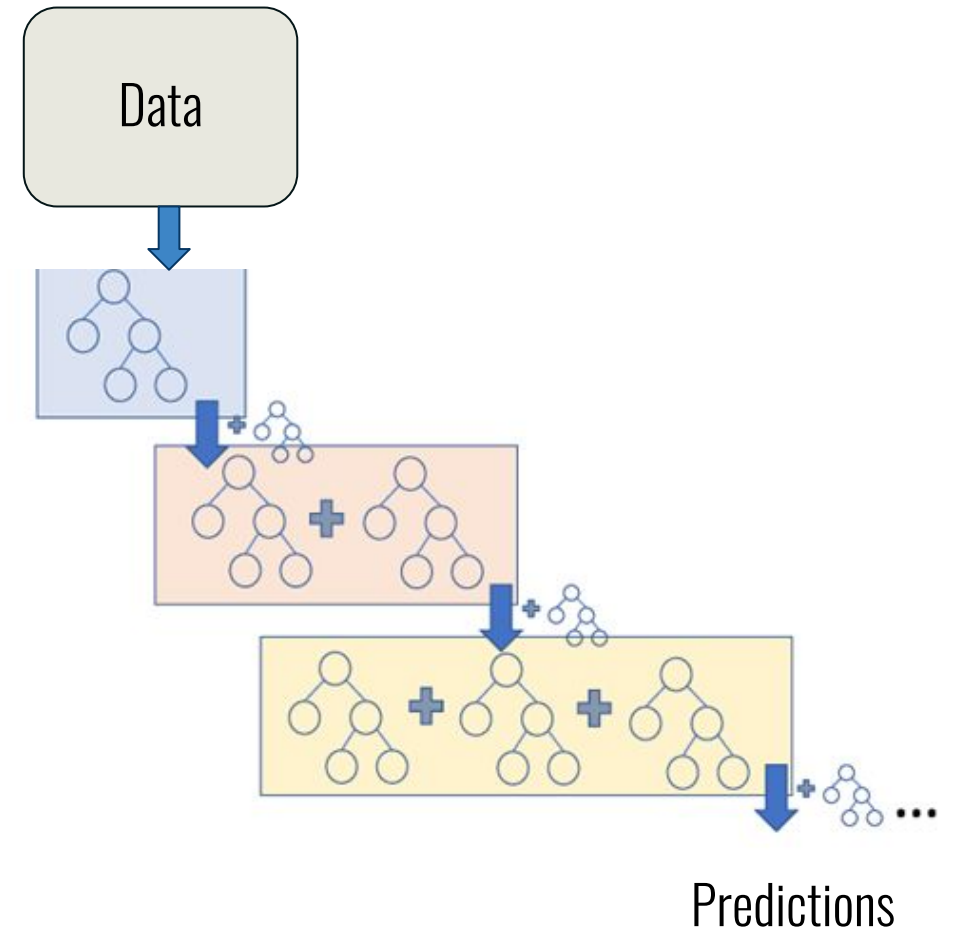
- **Ensemble learning**
Meta approach to modeling
Leverages the power of multiple predictive models
- **Random Forests**
Build all trees simultaneously and independently
- **Gradient Boosting**
Builds trees sequentially, with each tree being built in a chain.



Ensemble Model



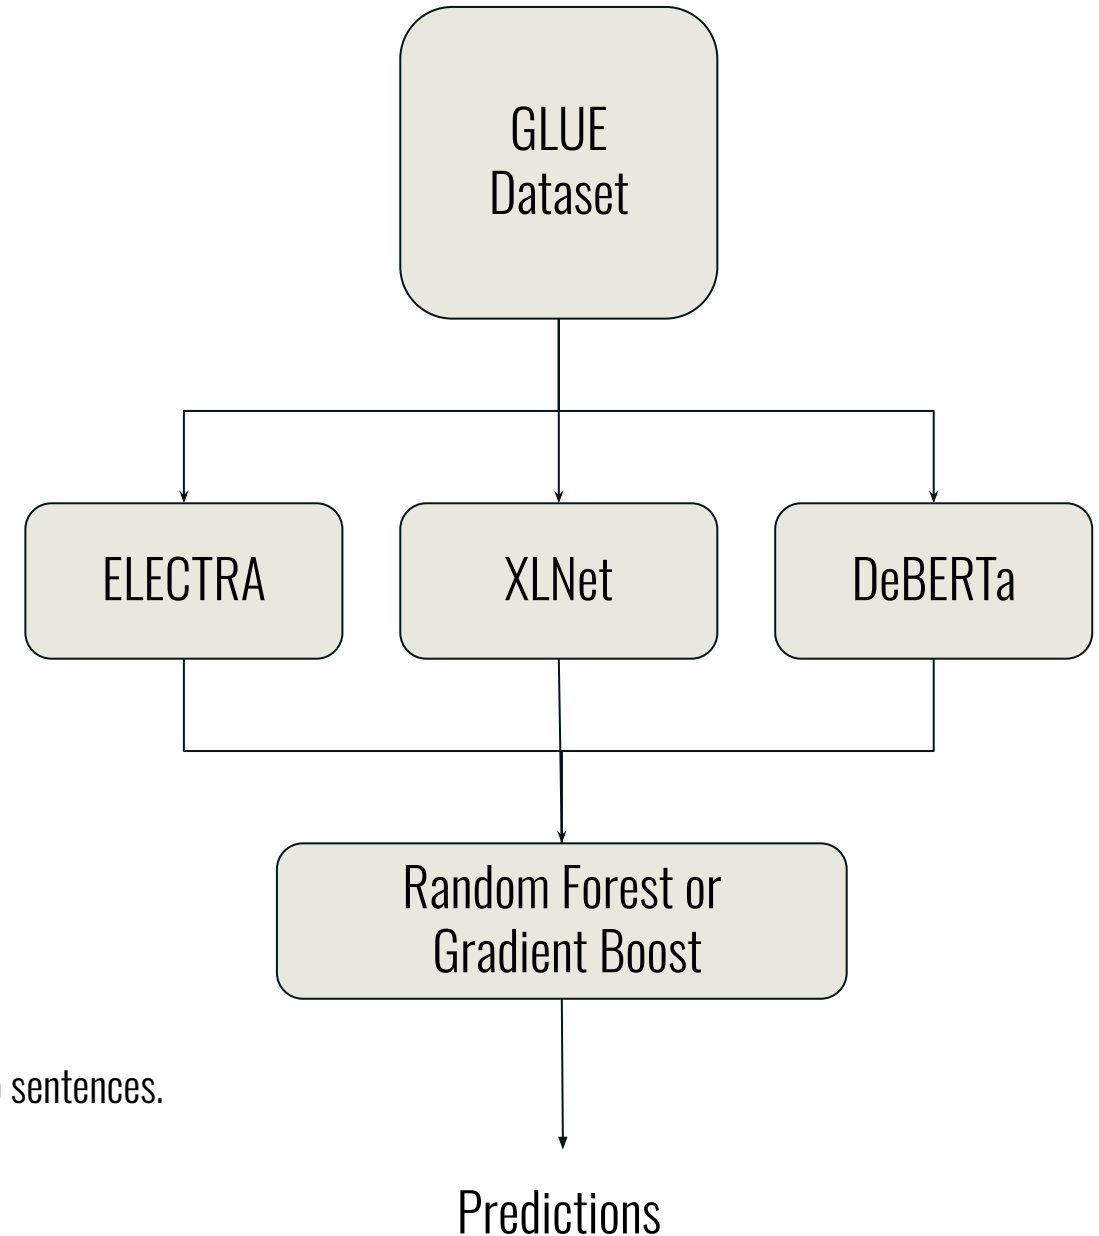
Random Forest Model



Gradient Boosting Model

Ensemble Model

- **Ensemble learning**
Meta approach to modeling
Leverages the power of multiple predictive models
- **Random Forests**
Build all trees simultaneously and independently (horizontally)
- **Gradient Boosting**
Builds trees sequentially, with each tree being built in a chain.
- **Classifier and Regressor Models**
STS-B has a continuous variable which measures the similarity of the two sentences.
All other tasks are binary or class based



Experiment Results

- Ensembles were competitive against each other
- Gradient Boosting outperformed, except on MRPC and RTE
- XLNet's long-term dependencies strong suite may have helped

		Transformers			Ensemble Methods	
Corpus	Metric	Electra	XLNet	Deberta	Random Forest	Grad Boost
Single-Sentence Tasks						
CoLA	Matthew's correlation	0.607	0.400	0.621	0.688	0.657
SST-2	Accuracy	0.917	0.940	0.935	0.958	0.954
Similarity and Paraphrase Tasks						
MRPC	Accuracy and F1	0.882 f1:0.915	0.892 f1:0.923	0.865 f1:0.903	0.878 f1:0.912	0.854 f1:0.893
QQP	Accuracy and F1	0.900 f1:0.866	0.874 f1:0.833	0.906 f1:0.875	0.909 f1:0.878	0.910 f1:0.879
STS-B	Pearson and Spearman	P:0.873 S:0.872	P:0.893 S:0.889	P:0.868 S:0.868	P:0.894 S:0.889 *Note: Regressor	P:0.897 S:0.892 *Note: Regressor
Inference Tasks						
MNLI	Accuracy	0.817	0.857	0.875	0.877	0.877
MNLI-MM	Accuracy	0.821	0.858	0.872	0.875	0.877
QNLI	Accuracy	0.889	0.878	0.915	0.916	0.924
RTE	Accuracy	0.682	0.740	0.668	0.726	0.702
WNLI	Accuracy	0.465	0.563	0.437	0.500	0.636

Conclusions

Results + Analysis

- Ensembling works!
- Head-to-head: DeBERTa and XLNet
- Gradient-Booster Ensemble was our winner

Limitations & Future Work

- 12-Hour VPN limit
- Employing TPUs for handling computationally expensive models
- Add different models such as T5 or LUKE to Ensemble



References

[GLUE Explained: Understanding BERT Through Benchmarks](#)

[Dataset description sheet](#)

[GLUE: A MULTI-TASK BENCHMARK AND ANALYSIS PLATFORM FOR NATURAL LANGUAGE UNDERSTANDING](#)

[DeBERTa Original Paper](#)

[Hugging Face Transformers Examples](#)

[ELECTRA: Pre-training Text Encoders as Discriminators Rather Than Generators.](#)

[Paper Reading #2: XLNet Explained](#)

[XLNet: Generalized Autoregressive Pretraining for Language Understanding](#)

[Guide to XLNet for Language Understanding](#)

[Data Science Central](#)

[XLNet Explained](#)

[**XLNet**: Generalized Autoregressive Pre-training for Language Understanding](#)

ELECTRA SMALL

Built with 12 layers, 256 hidden size and 14M parameters

```
(electra): ElectraModel(
  (embeddings): ElectraEmbeddings(
    (word_embeddings): Embedding(30522, 128, padding_idx=0)
    (position_embeddings): Embedding(512, 128)
    (token_type_embeddings): Embedding(2, 128)
    (LayerNorm): LayerNorm((128,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (embeddings_project): Linear(in_features=128, out_features=256, bias=True)
  (encoder): ElectraEncoder(
    (layer): ModuleList(
      (0): ElectraLayer(
        (attention): ElectraAttention(
          (self): ElectraSelfAttention(
            (query): Linear(in_features=256, out_features=256, bias=True)
            (key): Linear(in_features=256, out_features=256, bias=True)
            (value): Linear(in_features=256, out_features=256, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
          (output): ElectraSelfOutput(
            (dense): Linear(in_features=256, out_features=256, bias=True)
            (LayerNorm): LayerNorm((256,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
    )
    (intermediate): ElectraIntermediate(
      (dense): Linear(in_features=256, out_features=1024, bias=True)
    )
    (output): ElectraOutput(
      (dense): Linear(in_features=1024, out_features=256, bias=True)
      (LayerNorm): LayerNorm((256,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
)
```

```
( vocab_size = 30522, embedding_size = 128, hidden_size = 256,
num_hidden_layers = 12, num_attention_heads = 4, intermediate_size =
1024, hidden_act = 'gelu', hidden_dropout_prob = 0.1,
attention_probs_dropout_prob = 0.1, max_position_embeddings = 512,
type_vocab_size = 2, initializer_range = 0.02, layer_norm_eps = 1e-
12, summary_type = 'first', summary_use_proj = True,
summary_activation = 'gelu', summary_last_dropout = 0.1, pad_token_id
= 0, position_embedding_type = 'absolute', use_cache = True,
classifier_dropout = None, **kwargs )
```

```
(classifier): ElectraClassificationHead(
  (dense): Linear(in_features=256, out_features=256, bias=True)
  (dropout): Dropout(p=0.1, inplace=False)
  (out_proj): Linear(in_features=256, out_features=2, bias=True)
)
```